

1. Preprocessing Plugin

Purpose

To **prepare raw images** for deep-learning workflows. Offers both **single image preview** and **bulk** operations. Allows you to correct illumination, filter noise, extract channels, overlay separate channels, and more.

Major Sections

1.1 Single-Image Preprocessing

1. **Load Image for Preview**
 - **Browse Image...:** Opens a file dialog to select one image.
 - **Next/Prev Image:** Navigates through images in the same folder (if multiple exist).
 - **Original / Processed Preview Windows:** Shows side-by-side previews.
2. **Basic Options**
 - **Grayscale:** Convert the image from BGR to a single grayscale channel.
 - **Invert:** Invert pixel intensities (useful for certain microscopic images).
3. **Thresholding**
 - **None:** No threshold is applied.
 - **Manual:** Applies a user-specified threshold value.
 - **Adaptive:** Uses adaptive thresholding (per-channel for color images, or single-channel if grayscale).
 - **Otsu:** Computes the optimal threshold via Otsu's method (per-channel if color).
4. **Histogram Equalization & CLAHE**
 - **Histogram Equalization:** Globally enhances contrast by equalizing the pixel intensity histogram.
 - **CLAHE** (Contrast Limited Adaptive Histogram Equalization): More localized enhancement than global hist-eq, especially beneficial for medical images.
5. **Morphological Operations**
 - **Operation:** Erode, Dilate, Opening, Closing, TopHat, BlackHat.
 - **Kernel Size:** Size of the structuring element (e.g., 3×3).
6. **Denoising**
 - **Gaussian / Median / Bilateral / NLM:** Various noise reduction filters.
 - **Kernel:** Controls the filter's neighborhood size or parameters.
7. **Channel Extraction**
 - **B, G, R:** Optionally extract a single color channel if needed.
 - If "None," no specific channel extraction is performed.
8. **Resize**
 - **Width & Height:** If non-zero, resizes images to those dimensions.
9. **Enhance Red Channel**

- **Enable:** Toggles whether to apply brightness/contrast specifically to the red channel.
- **Brightness Factor & Contrast Factor:** Fine-tunes how the red channel is brightened or how its contrast is amplified.

10. Apply Preprocessing

- Executes all chosen transformations in the specified order.
- The result is shown in the **Processed** preview.

11. Save Processed Image

- Exports the single processed image to disk, with any extension (.png, .jpg, .tiff, etc.).

1.2 Bulk Preprocessing

1. **Input/Output Folders**
 - **Input:** Folder containing images to preprocess.
 - **Output:** Destination folder to save processed images.
2. **Output Format & Suffix**
 - **Convert to Format:** jpg, png, or tiff.
 - **Output Suffix:** e.g., _processed. Final filenames might look like myimage_processed.jpg.
3. **Process Bulk Images**
 - Runs a **QThread** so the UI remains responsive.
 - Applies the same pipeline configured in single-image mode (i.e., the same thresholding, morphological ops, etc.) to all images under `Input` recursively.

1.3 Overlay (ch01 -> Red, ch02 -> Green)

For combining multiple monochrome channel images into a single color-coded overlay:

1. **Input/Output:** Similar to bulk preprocessing, specify in/out folders.
2. **ch01 substring / ch02 substring:** Identifiers in the filenames that separate the “red channel” images (-ch01) from the “green channel” images (-ch02).
3. **Normalize Channels:** Each channel is scaled from min=0 to max=255.
4. **Hist Equal:** Global histogram equalization before merging.
5. **CLAHE:** Local histogram enhancement before merging.
6. **Overlay Suffix:** e.g., _overlay, added to the final merged file.

1.4 Filter by Substring

1. **Input/Output:** Folders for the filter operation.
2. **Substring to Keep:** e.g., “ch01”. Only files containing this substring in the name are copied.

3. **Filter and Copy Images:** Useful for quickly splitting channels or partial subsets from a large data folder.
-

2. Data Exploration Plugin

Purpose

To **explore dataset structure** before training. Generates bar plots of class distributions, logs data counts, and displays sample images per class.

Major Features

1. **Data Folder**
 - Choose the parent folder containing subfolders named after each class (e.g., data/ClassA, data/ClassB, etc.).
2. **Max Images/Class**
 - Limits the number of sample images displayed in the gallery for each class (default: 3).
3. **Analyze Dataset**
 - **Scans each subfolder** to count images (jpg, png, tif, etc.).
 - **Displays:**
 - **Bar Chart:** Number of images for each class.
 - **Gallery:** A scrollable preview of sample images.
 - **Logs:** If any classes have zero images or if the folder structure is malformed.
4. **Progress Bar & Log**
 - The scanning runs in a background thread. The log displays real-time updates, such as “Class 'A': 120 images found.”

Useful for verifying you have enough images per class and quickly spotting data imbalance issues.

3. Training Plugin

Purpose

To **train convolutional neural networks** with a wide variety of parameters and advanced features. Uses **PyTorch Lightning** under the hood.

Key Configuration Areas

3.1 Dataset Splits

1. **Dataset Folder:** Root directory containing subfolders for each class.
2. **Val Split & Test Split** (percentages): The plugin will split your dataset accordingly.
 - E.g., 15% for validation, 15% for test → 70% training.

3.2 Network Architecture

Choose from:

- resnet18, resnet50, resnet101
- densenet
- vgg
- inception
- mobilenet
- efficientnet_b0
- convnext_tiny
- convnext_large

Toggle **Use 299 for Inception** if using the Inception architecture, to handle 299×299 input.

3.3 Augmentations

- **Brightness/Contrast:** Randomly adjusts image brightness and contrast.
- **Hue/Saturation:** Perturbs color balance.
- **Gauss Noise:** Injects random Gaussian noise.
- **Rotation:** Up to $\pm 30^\circ$ (configurable) rotation.
- **Flipping:** Horizontal, vertical, or both. Also specify flip probability.
- **Random Crop:** Randomly crops and resizes images.
- **Elastic:** Elastic transformations for warping images.
- **Grid Distortion:** Distorts image via a grid warping.
- **Optical Distortion:** Simulated lens distortion.
- **Normalize Pixel:** Typical ImageNet mean/std normalization.
- **MixUp / CutMix:**
 - Blends or cuts-and-pastes pairs of images with a Beta(α , α) sampling to augment training.

3.4 Hyperparameters

- **Learning Rate:** Initial LR, e.g., 1e-4.
- **Momentum:** For SGD-based optimizers.
- **Weight Decay:** L2 regularization.
- **Optimizer:** adam, sgd, adamw, or lamb (if installed).
- **Scheduler:** none, steplr, reducelronplateau, cosineannealing, cyclicalr.
- **Scheduler Params:** Additional parameters, e.g., step_size=10, gamma=0.1.

3.5 Training Controls

- **Epochs:** Maximum number of epochs to train.
- **Batch Size:** e.g., 8, 16, 32...
- **Early Stopping:** Halts training if validation metric does not improve for N epochs.
 - Additional fields: **Monitor Metric** (`val_loss` or `val_acc`), **Patience**, **Min Delta**, **Mode** (`min/max`).

3.6 Regularization and Advanced Settings

- **Dropout Rate:** Probability of dropout before the final classifier.
- **Label Smoothing:** Softens hard labels in cross-entropy.
- **Freeze Entire Backbone:** Disables gradient updates in the backbone (fine-tuning only the final layer).
- **Loss Function:**
 - `cross_entropy`, `focal`, `bce` (for single-class but still multi-output?), `bce_single_logit` (strictly 1-class).
- **Gradient Clip Val:** Prevent exploding gradients.
- **LR Finder:** Automatic learning-rate suggestion.
 - **Accept LR Suggestion:** If checked, applies the found LR to training.
- **TensorBoard Logger:** Logs to `tb_logs/` for visualization in TensorBoard.
- **Mixed Precision:** If checked, uses `precision=16` (FP16) for faster GPU training.
- **Warmup Epochs:** Allows a warmup schedule before the main LR policy.
- **Workers:** Number of multiprocessing workers to load data.
- **Gradient Checkpointing:** Saves memory by trading off some compute overhead.
- **Gradient Accumulation:** Accumulate gradients over N mini-batches before stepping.
- **Check Val Every N Epoch:** Frequency of validation.

3.7 Partial Freezing Options (ResNet/ConvNeXt)

- For ResNet-like architectures: `conv1_bn1`, `layer1`, `layer2`, `layer3`, `layer4`.
- For ConvNeXt: `convnext_block0..convnext_block3`.
Allows selective freezing of certain stages.

3.8 Training Execution

- **Start Training:** Kicks off a background thread. Real-time logs and progress updates appear in the text field.
- After training completes, a “best checkpoint” (`.ckpt`) is saved and test set metrics (confusion matrix, classification report) are displayed.

3.9 Optuna Hyperparameter Tuning (Optional)

- **Tune with Optuna:** Automates search for best LR, dropout, etc.
- **Trials:** Number of runs.
- **Timeout:** Maximum seconds to spend searching (0 = no limit).
- **Use Test Loss as Optuna Objective:** If checked, will optimize on the test metric (not recommended in typical practice).

3.10 Export Results

- After training, export confusion matrix, classification report, and optionally zip up TensorBoard logs.

4. Evaluation/Inference Plugin

Purpose

To **deploy or test** a trained model on images (single or batch). Allows Grad-CAM overlays, misclassification analysis, and advanced interpretability (SHAP).

Main Features

4.1 Checkpoint Selection

- **Browse Checkpoint...:** Load a `.ckpt` from the training stage.
- Automatically parses model architecture, class names, etc.

4.2 Transform Customization

- Similar to training transforms: e.g., resize, normalization means, standard deviations.

4.3 Single Image Inference

1. **Test Image:** Choose one image.
2. **Run Inference (Single):** Predicts classes for that image.
 - **Top-k:** Show top-k predictions and their confidence scores.
 - **Min Confidence:** Filter out predictions below this percentage.
3. **Grad-CAM++:** If checked, calculates and overlays a heatmap (Grad-CAM or Grad-CAM++) highlighting regions most significant for the predicted class.

4.4 Evaluation & Batch Inference

1. **Inference Folder:** A directory of images to predict in bulk.
2. **GT CSV (optional):** If you have ground-truth labels (filename, label) in a CSV, the plugin calculates metrics: confusion matrix, classification report, etc.

3. **Batch Grad-CAM:** For each image, generate a Grad-CAM overlay.
4. **Export CSV:** Save a `predictions.csv` with top-k predictions for each file.
Also, `misclassified.csv` if ground truth is provided.
5. **SHAP Analysis** (if `shap` is installed):
 - **Enable SHAP:** Produces a SHAP summary plot for a subset of images.
 - **SHAP Samples:** Number of images to visualize via SHAP.
 - **SHAP BG Samples:** Size of the background set.

4.5 Final Outputs:

- **Confusion Matrix:** For classification results, saved to text and PNG.
- **ROC/PR Curves:** If binary or multi-class.
- **Calibration Curve:** For binary.
- **Classification Report:** Precision, recall, F1.
- **Metrics JSON:** Summaries stored in `metrics.json`.